

```
In [206... '''
Patrick Ballou
ID: 801130521
ECGR 4105
Homework 2
Problem 1
'''
```

```
Out[206]: '\nPatrick Ballou\nID: 801130521\nECGR 4105\nHomework 2\nProblem 1\n'
```

```
In [207... import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
In [208... df = pd.read_csv("diabetes.csv")
df.head()
```

```
Out[208]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

```
In [209... #train/test split, random_state functions as seed
df_train, df_test = train_test_split(df, train_size=.8, test_size=.2, random_state=7)
```

```
In [210... #separate dataset into x and y, train and test
inputs = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

x_train = df_train[inputs]
Y_train = df_train['Outcome']

x_test = df_test[inputs]
Y_test = df_test['Outcome']
```

```
In [211... #standard scaler performs best
scaler = StandardScaler()
#scaler = MinMaxScaler()
X_train = scaler.fit_transform(x_train)
X_test = scaler.fit_transform(x_test)
```

```
In [212... #for loop finds best C value, which is the default in this case
'''
C = [10, 1, .1, .01, .001]
```

```

for c in C:
    classifier = LogisticRegression(random_state=7, C=c)
    classifier.fit(X_train, Y_train)
    print("C:", c)
    print("Training accuracy:", classifier.score(X_train, Y_train))
    print("Testing accuracy:", classifier.score(X_test, Y_test))

...

classifier = LogisticRegression(random_state=7)
classifier.fit(X_train, Y_train)

```

Out[212]:

```

▼      LogisticRegression
LogisticRegression(random_state=7)

```

In [213... `Y_pred = classifier.predict(X_test)`In [214... `cnf_matrix = metrics.confusion_matrix(Y_test, Y_pred)`
`cnf_matrix`Out[214]: `array([[89, 8],
 [25, 32]], dtype=int64)`In [215... `print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred))`
`print("Precision:", metrics.precision_score(Y_test, Y_pred))`
`print("Recall:", metrics.recall_score(Y_test, Y_pred))`

Accuracy: 0.7857142857142857

Precision: 0.8

Recall: 0.5614035087719298

In [216... *#not a great classifier*
`cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix=cnf_matrix, display_labels=)`
`cm_display.plot()`Out[216]: `<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1cf4c777df0>`